

# CSC 555 Mining Big Data

## Project, Phase 1

Lavinia Wang #1473704

In this part of the project (which will also serve as our take-home midterm), you will 1) Set up a 3un-node cluster and 2) perform data warehousing and transformation queries using Hive, Pig and Hadoop streaming. The modified Hive-style schema is at:

[http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/SSBM\\_schema\\_hive.sql](http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/SSBM_schema_hive.sql)

It is based on SSBM benchmark (derived from industry standard TPCB benchmark). The data is at Scale1, or the smallest unit – lineorder is the largest table at about 0.6GB. You can use wget to download the following links. Keep in mind that data is | -separated (not CSV).

<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/dwdate.tbl>

<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/lineorder.tbl>

<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/part.tbl>

<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/supplier.tbl>

<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/customer.tbl>

Please be sure to submit all code (pig, python and SQL).

## Part 1: Multi-node cluster

- 1) Your first step is to setup a multi-node cluster and re-run a simple wordcount. For this part, you will create a 3-node cluster (with a total of 1 master + 2 worker nodes). Include your master node in the “slaves” file, to make sure **all 3** nodes are working.

You need to perform the following steps:

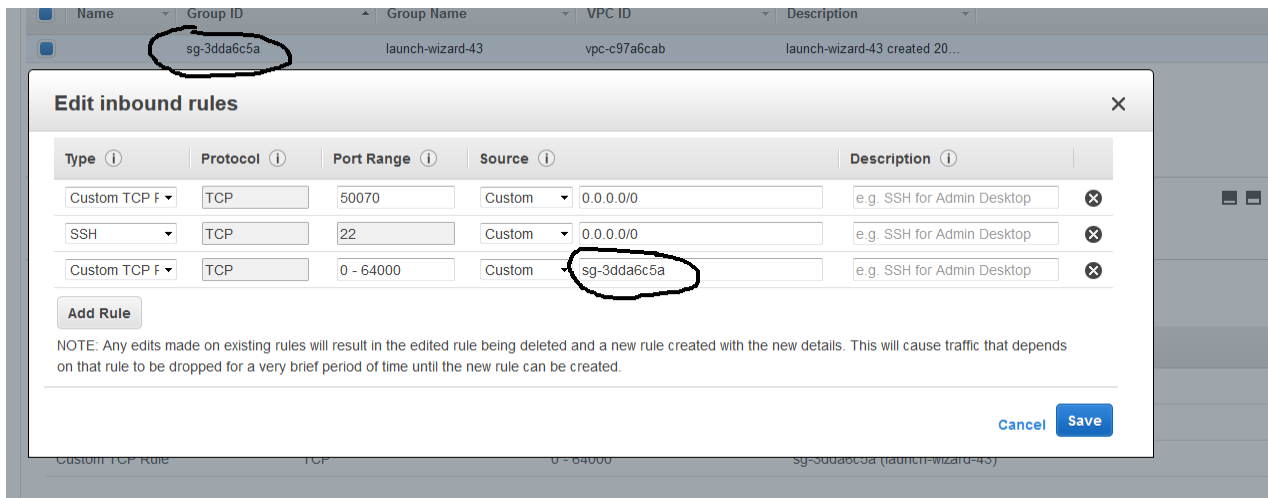
1. Create a new node of a medium size (you can always switch the size of the node). It is possible, but I do not recommend trying to reconfigure your existing Hadoop into this new cluster (it is much easier to make 3 new nodes for a total of 4 in your AWS account).
  - a. **When creating a node I recommend changing the default 8G hard drive to 30G on all nodes.**
  - b. Change your security group setting to open firewall access. We need to open the ports in two different ways. We will open port 50070 for the web interface in order to be able to see the cluster status in a browser. We will also set 0-64000 range opening up all ports. However, we will ensure that the ports are open only **within** the cluster and not to the world.

In order to make changes, you need to do the following. Access the cluster security group (launch-wizard-xx).

Elastic IPs	
Availability zone	us-west-1b
Security groups	launch-wizard-39. <a href="#">view rules</a>
Scheduled events	-

Right click on the security group and choose Edit inbound rules

Note that the first line below is opening port 50070. The second line below is the default (port 22 is required for regular SSH connections). The third line opens all ports but ONLY for the same security group (assuming that all of your nodes in the cluster share the same security group – that will happen automatically if you use the “create more like this” option when creating instances as specified in part 1-c below). We previously had some issues with machines being hacked without that last limitation, so make sure you include it.



- c. Right click on the Master node and choose “create more like this” to create 2 more nodes with same settings. If you configure the network settings on master first, security group information will be copied.

**NOTE: Hard drive size will not be copied and default to 8G unless you change it.**

2. Connect to the master and set up Hadoop similarly to what you did previously. Do not attempt to repeat these steps on workers yet – you will only need to set up Hadoop once.

- a. Configure core-site.xml, adding the **PrivateIP** (do not use public IP) of the master.

```

limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>

<property>
<name>fs.defaultFS</name>
<value>hdfs://172.31.7.201/</value>
</property>

</configuration>
[ec2-user@ip-172-31-7-201 ~]$ cat hadoop-2.6.4/etc/hadoop/core-site.xml

```

- b. Configure hdfs-site and set replication factor to 2.

```
<!-- Put site-specific property overrides in this file. -->

<configuration>

<property>
<name>dfs.replication</name>
<value>2</value>
</property>

</configuration>
[ec2-user@ip-172-31-9-105 ~]$
```

- c. cp `hadoop-2.6.4/etc/hadoop/mapred-site.xml.template` `hadoop-2.6.4/etc/hadoop/mapred-site.xml` and then configure `mapred-site.xml`

```
<!-- Put site-specific property overrides in this file. -->

<configuration>

<property>
<name>mapreduce.framework.name</name>
<value>yarn</value>
</property>

</configuration>
[ec2-user@ip-172-31-9-105 ~]$ cat hadoop-2.6.4/etc/hadoop/mapred-site.xml
```

- d. Configure `yarn-site.xml` (once again, use PrivateIP of the master)

```
<!-- Site specific YARN configuration properties -->

<property>
<name>yarn.resourcemanager.hostname</name>
<value>172.31.7.201</value>
</property>

<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>

</configuration>
[ec2-user@ip-172-31-7-201 ~]$ cat hadoop-2.6.4/etc/hadoop/yarn-site.xml
```

Finally, edit the `slaves` file and list your 4 nodes (master and 3 workers) using Private IPs

```
[ec2-user@ip-172-31-7-201 ~]$ cat hadoop-2.6.4/etc/hadoop/slaves
172.31.7.201
172.31.5.246
...
```

Make sure that you use private IP (private DNS is also ok) for your configuration files (such as `conf/masters` and `conf/slaves` or the other 3 config files). The advantage of the Private IP is that it does not change after your instance is stopped (if you use the Public IP, the cluster would need to be reconfigured every time it is stopped). The downside of the Private IP is that it is only meaningful within the Amazon EC2 network. So all nodes in EC2 can talk to each other using

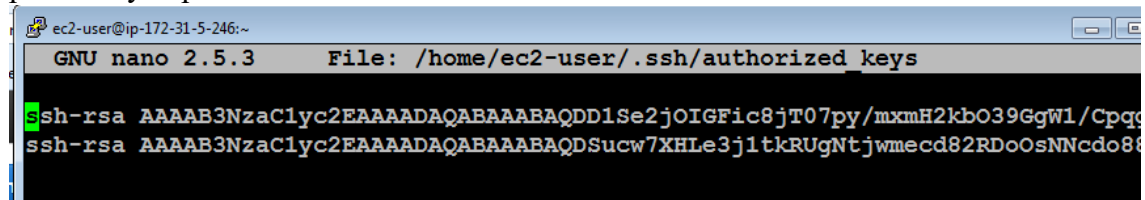
Private IP, but you cannot connect to your instance from the outside (e.g., from your laptop) because Private IP has no meaning for your laptop (since your laptop is not part of the Amazon EC2 network).

Now, we will pack up and move Hadoop to the workers. All you need to do is to generate and then copy the public key to the worker nodes to achieve passwordless access across your cluster.

1. Run `ssh-keygen -t rsa` (and enter empty values for the passphrase) on the master node. That will generate `.ssh/id_rsa` and `.ssh/id_rsa.pub` (private and public key). You now need to manually copy the `.ssh/id_rsa.pub` and append it to `~/.ssh/authorized_keys` **on each worker.**

Keep in mind that this is a single-line public key and accidentally introducing a line break would cause a mismatch.

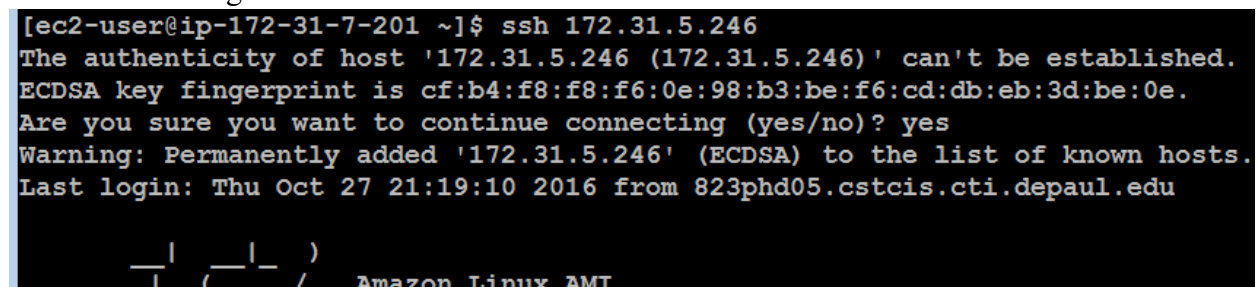
Note that the example below is NOT the master, but one of the workers (ip-172-31-5-246). The first public key is the .pem Amazon half and the 2<sup>nd</sup> public key is the master's public key copied in as one line.



```
ec2-user@ip-172-31-5-246:~  
GNU nano 2.5.3 File: /home/ec2-user/.ssh/authorized_keys  
ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQDD1Se2jOIGFic8jT07py/mxmH2kbO39GgW1/Cpqq  
ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQDSucw7XHL3j1tkRUgNtjwmecd82RDoOsNNcdo88
```

You can add the public key of the master to the master by running this command:  
`cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys`

Make sure that you can ssh to all of the nodes from the master node (by running ssh 54.186.221.92, where the IP address is your worker node) from the master and ensuring that you were able to login. You can exit after successful ssh connection by typing `exit` (the command prompt will tell you which machine you are connected to, e.g., `ec2-user@ip-172-31-37-113`). Here's me ssh-ing from master to worker.



```
[ec2-user@ip-172-31-7-201 ~]$ ssh 172.31.5.246  
The authenticity of host '172.31.5.246 (172.31.5.246)' can't be established.  
ECDSA key fingerprint is cf:b4:f8:f8:f6:0e:98:b3:be:f6:cd:db:eb:3d:be:0e.  
Are you sure you want to continue connecting (yes/no)? yes  
Warning: Permanently added '172.31.5.246' (ECDSA) to the list of known hosts.  
Last login: Thu Oct 27 21:19:10 2016 from 823phd05.cstcis.cti.depaul.edu  
  
_ | _ | )  
| ( _ | / Amazon Linux AMI
```

Once you have verified that you can ssh from the master node to every cluster member including the master itself (ssh localhost), you are going to return to the master node (**exit** until your prompt shows the IP address of the master node) and pack the contents of the hadoop directory there. Make sure your Hadoop installation is configured correctly (because from now on, you will have 4 copies of the Hadoop directory and all changes need to be applied in 4 places).

`cd` (go to root home directory, i.e. `/home/ec2-user/`)

(pack up the entire Hadoop directory into a single file for transfer. You can optionally compress the file with gzip)

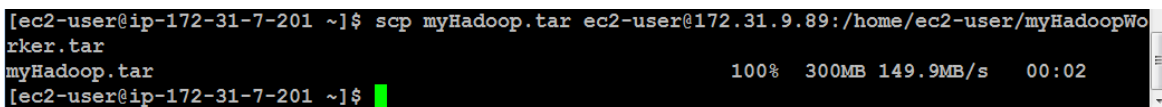
**tar cvf myHadoop.tar hadoop-2.6.4**

**ls -al myHadoop.tar** (to verify that the .tar file had been created)

Now, you need to copy the myHadoop.tar file to every non-master node in the cluster. If you had successfully setup public-private key access in the previous step, this command (for each worker node) will do that:

(copies the myHadoop.tar file from the current node to a remote node into a file called myHadoopWorker.tar. Don't forget to replace the IP address with that your worker nodes. By the way, since you are on the Amazon EC2 network, either Public or Private IP will work just fine.)

**scp myHadoop.tar ec2-user@54.187.63.189:/home/ec2-user/myHadoopWorker.tar**



```
[ec2-user@ip-172-31-7-201 ~]$ scp myHadoop.tar ec2-user@172.31.9.89:/home/ec2-user/myHadoopWorker.tar
myHadoop.tar                                100% 300MB 149.9MB/s   00:02
[ec2-user@ip-172-31-7-201 ~]$
```

Once the tar file containing your Hadoop installation from master node has been copied to each worker node, you need to login to each non-master node and unpack the .tar file.

Run the following command (on each worker node, not on the master) to untar the hadoop file. We are purposely using a different tar archive name (i.e., **myHadoopWorker.tar**), so if you get “file not found” error, that means you are running this command on the master node or have not yet successfully copied myHadoopWorker.tar file to the worker.

**tar xvf myHadoopWorker.tar**

Once you are done, run this on the master (nothing needs to be done on the workers to format the cluster unless you are re-formatting, in which case you'll need to delete the dfs directory).

**hadoop namenode -format**

Once you have successfully completed the previous steps, you should can start and use your new cluster by going to the master node and running the start-dfs.sh and start-yarn.sh scripts (you do not need to explicitly start anything on worker nodes – the master will do that for you).

You should verify that the cluster is running by pointing your browser to the link below.

[http://\[insert-the-public-ip-of-master\]:50070/](http://[insert-the-public-ip-of-master]:50070/)

Make sure that the cluster is operational (you can see the 4 nodes under Datanodes tab).

Submit a screenshot of your cluster status view.

## Datanode Information

### In operation

Node	Last contact	Admin State	Capacity	Used	Non DFS Used	Remaining	Blocks	Block pool used	Failed Volumes	Version
ip-172-31-6-219.us-east-2.compute.internal (172.31.6.219:50010)	1	In Service	29.99 GB	4 KB	2.03 GB	27.96 GB	0	4 KB (0%)	0	2.6.4
ip-172-31-9-235.us-east-2.compute.internal (172.31.9.235:50010)	1	In Service	29.99 GB	8 KB	2.54 GB	27.45 GB	0	8 KB (0%)	0	2.6.4
ip-172-31-7-186.us-east-2.compute.internal (172.31.7.186:50010)	1	In Service	29.99 GB	4 KB	2.02 GB	27.96 GB	0	4 KB (0%)	0	2.6.4

Repeat the steps for wordcount using bioproject.xml from Assignment 1 and submit screenshots of running it.

```

Downloads — ec2-user@ip-172-31-9-235:~ — ssh -i lwang73.pem ec2-user@ec2-3-16-76-250.us-east-2.com..
19/05/20 05:11:31 INFO client.RMProxy: Connecting to ResourceManager at /172.31.9.235:8032
19/05/20 05:11:32 INFO input.FileInputFormat: Total input paths to process : 1
19/05/20 05:11:32 INFO mapreduce.JobSubmitter: number of splits:2
19/05/20 05:11:32 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1558322887419_0020
19/05/20 05:11:32 INFO impl.YarnClientImpl: Submitted application application_1558322887419_0020
19/05/20 05:11:32 INFO mapreduce.Job: The url to track the job: http://ip-172-31-9-235.us-east-2.compute.i
nternal:8088/proxy/application_1558322887419_0020/
19/05/20 05:11:32 INFO mapreduce.Job: Running job: job_1558322887419_0020
19/05/20 05:11:39 INFO mapreduce.Job: Job job_1558322887419_0020 running in uber mode : false
19/05/20 05:11:39 INFO mapreduce.Job:  map 0% reduce 0%
19/05/20 05:11:50 INFO mapreduce.Job:  map 26% reduce 0%
19/05/20 05:11:53 INFO mapreduce.Job:  map 46% reduce 0%
19/05/20 05:11:56 INFO mapreduce.Job:  map 48% reduce 0%
19/05/20 05:11:59 INFO mapreduce.Job:  map 77% reduce 0%
19/05/20 05:12:02 INFO mapreduce.Job:  map 83% reduce 0%
19/05/20 05:12:06 INFO mapreduce.Job:  map 100% reduce 0%
19/05/20 05:12:09 INFO mapreduce.Job:  map 100% reduce 100%
19/05/20 05:12:09 INFO mapreduce.Job: Job job_1558322887419_0020 completed successfully
19/05/20 05:12:09 INFO mapreduce.Job: Counters: 49
    File System Counters
      FILE: Number of bytes read=59605201
      FILE: Number of bytes written=86827979

```

## Lavinia Wang #1473704

```
Downloads — ec2-user@ip-172-31-9-235:~ — ssh -i lwang73.pem ec2-user@ec2-3-16-76-250.us-east-2.com..
FILE: Number of bytes written=86827979
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=231153307
HDFS: Number of bytes written=20056175
HDFS: Number of read operations=9
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
Job Counters
  Launched map tasks=2
  Launched reduce tasks=1
  Data-local map tasks=2
  Total time spent by all maps in occupied slots (ms)=44447
  Total time spent by all reduces in occupied slots (ms)=6333
  Total time spent by all map tasks (ms)=44447
  Total time spent by all reduce tasks (ms)=6333
  Total vcore-milliseconds taken by all map tasks=44447
  Total vcore-milliseconds taken by all reduce tasks=6333
  Total megabyte-milliseconds taken by all map tasks=45513728
  Total megabyte-milliseconds taken by all reduce tasks=6484992
Map-Reduce Framework
Map-Reduce Framework
  Map input records=5284546
  Map output records=18562366
  Map output bytes=279356680
  Map output materialized bytes=26902454
  Input split bytes=208
  Combine input records=20053191
  Combine output records=2673165
  Reduce input groups=1040390
  Reduce shuffle bytes=26902454
  Reduce input records=1182340
  Reduce output records=1040390
  Spilled Records=3855505
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=569
  CPU time spent (ms)=39990
  Physical memory (bytes) snapshot=745447424
  Virtual memory (bytes) snapshot=6448369664
  Total committed heap usage (bytes)=526909440
Shuffle Errors
```



```
Downloads — ec2-user@ip-172-31-9-235:~ — ssh -i lwang73.pem ec2-user@ec2-3-16-76-250.us-east-2.com..
Merged Map outputs=2
GC time elapsed (ms)=569
CPU time spent (ms)=39990
Physical memory (bytes) snapshot=745447424
Virtual memory (bytes) snapshot=6448369664
Total committed heap usage (bytes)=526909440

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
Bytes Read=231153099
File Output Format Counters
Bytes Written=20056175

real    0m39.222s
user    0m3.779s
sys     0m0.263s
[ec2-user@ip-172-31-9-235 ~]$
```

```
[ec2-user@ip-172-31-9-235 ~]$ hadoop fs -du /data/wordcount1/
0          /data/wordcount1/_SUCCESS
20056175   /data/wordcount1/part-r-00000
```

```
Downloads — ec2-user@ip-172-31-9-235:~ — ssh -i lwang73.pem ec2-user@ec2-3-16-76-250.us-east-2.com..
[[ec2-user@ip-172-31-9-235 ~]$ hadoop fs -cat /data/wordcount1/part-r-00000 | grep arctic
<I>holarctica</I> 28
<I>holarctica</I></B>. 8
<I>holarctica</I>, 1
<I>palearctica</I> 4
<I>holarctica</I> 1
(Antarctic 3
(Antarctica) 1
(Antarctica), 11
<Label>Antarctic 1
<Name>Antarctic 3
<Name>Antarctica 1
<Strain>Antarctic 1
<Title>Antarctic 5
Antarctic 137
Antarctic, 1
Antarctic. 2
Antarctic.</Description> 1
Antarctic.</Title> 1
Antarctic</Title> 4
Antarctica 16
Antarctica)</Title> 1
```



Downloads — ec2-user@ip-172-31-9-235:~ — ssh -i lwang73.pem ec2-user@ec2-3-16-76-250.us-east-2.com..

```

antarcticum</OrganismName>      3
antarcticus      31
antarcticus<i>    4
antarcticus<i><b>.      1
antarcticus).      1
antarcticus,      1
antarcticus</Name>      5
antarcticus</OrganismName>      5
arctic      21
arctica      27
arctica<i>    2
arctica<i>    3
arctica<i>    1
arctica.</Description>      2
arctica</Name>      5
arctica</OrganismName>      5
arcticus      31
arcticus<i>    2
arcticus</Name>      4
arcticus</OrganismName>      4
holarctica      77
humans.Antarctic      1

```

Downloads — ec2-user@ip-172-31-9-235:~ — ssh -i lwang73.pem ec2-user@ec2-3-16-76-250.us-east-2.com..

```

Antarctica)</Title>      1
Antarctica,      9
Antarctica.      24
Antarctica.&#x0D;      3
Antarctica.</Description>      19
Antarctica</Description>      2
Antarctica</Name>      1
Antarctica</Title>      6
Palearctic      1
Project">Antarctic      1
Subarctic      11
abbr="Antarctic      1
antarctic      5
antarctica      17
antarctica<i><b>.&#x0D;      2
antarctica,      4
antarctica</Name>      10
antarctica</OrganismName>      11
antarctica</Title>      1
antarcticum      32
antarcticum</Name>      3
antarcticum</OrganismName>      3

```

```
Downloads — ec2-user@ip-172-31-9-235:~ — ssh -i lwang73.pem ec2-user@ec2-3-16-76-250.us-east-2.com.
arctica</i> 2
arctica</i> 3
arctica</i>, 1
arctica.</Description> 2
arctica</Name> 5
arctica</OrganismName> 5
arcticus 31
arcticus</i> 2
arcticus</Name> 4
arcticus</OrganismName> 4
holarctica 77
humans.Antarctic 1
palearctica 66
palearctica</Name> 1
sub-Antarctic 4
sub-arctic 4
subantarctic 1
subantarcticus 7
subantarcticus</Name> 1
subantarcticus</OrganismName> 1
subarctic 21
```

Submit a short paragraph with a discussion about how the results compare (faster? slower? How much faster/slower? Due to what?)

The single-node Hadoop instance from assignment 1 ran the word count job in 1 minute 14.366 seconds. The three-node instance ran the job 39.222 seconds which is roughly twice as fast. The speed increase is due to the additional nodes, however, I would have expected a more than doubling of speed, based just on the number of nodes. It's likely that network speed and block distribution have impacted speed.

Running the following command shows that the file was only split into two blocks and where they are located.

```
[ec2-user@ip-172-31-9-235 ~]$ hdfs fsck /data/bioproject.xml -files -blocks
-llocations
Connecting to namenode via http://ip-172-31-9-235.us-east-2.compute.internal:50070
FSCK started by ec2-user (auth:SIMPLE) from /172.31.9.235 for path
/data/bioproject.xml at Mon May 20 05:26:12 UTC 2019
/data/bioproject.xml 231149003 bytes, 2 block(s): OK
0. BP-177962531-172.31.9.235-1558313129568:blk_1073741826_1002 len=134217728
repl=2 [172.31.7.186:50010, 172.31.6.219:50010]
1. BP-177962531-172.31.9.235-1558313129568:blk_1073741827_1003 len=96931275
repl=2 [172.31.7.186:50010, 172.31.6.219:50010]
```

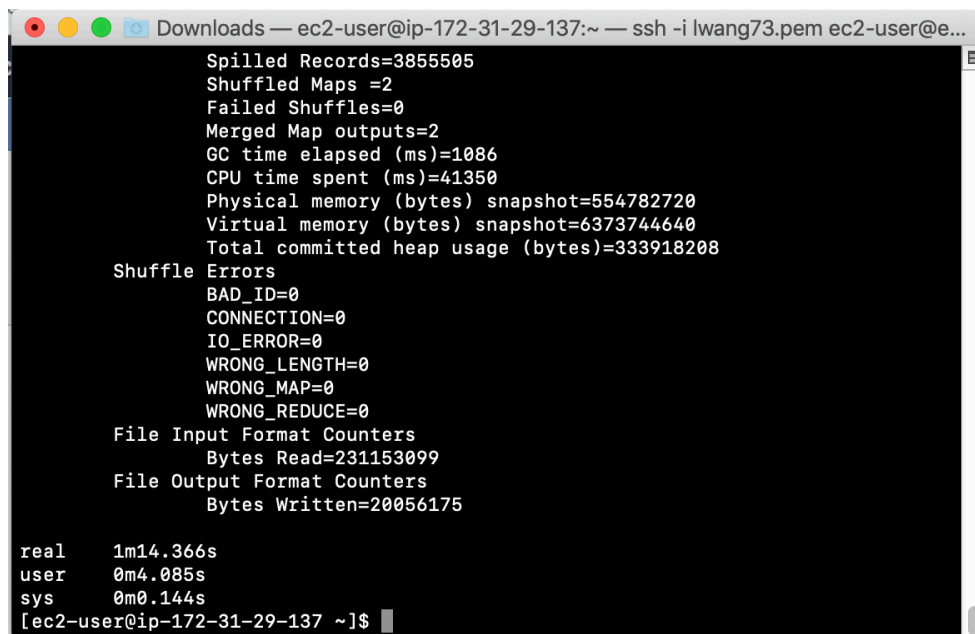
```
Status: HEALTHY
Total size: 231149003 B
Total dirs: 0
Total files: 1
Total symlinks: 0
Total blocks (validated): 2 (avg. block size 115574501 B)
Minimally replicated blocks: 2 (100.0 %)
Over-replicated blocks: 0 (0.0 %)
Under-replicated blocks: 0 (0.0 %)
```

```
Mis-replicated blocks:      0 (0.0 %)
Default replication factor: 2
Average block replication:   2.0
Corrupt blocks:             0
Missing replicas:           0 (0.0 %)
Number of data-nodes:       3
Number of racks:            1
FSCK ended at Mon May 20 05:26:12 UTC 2019 in 2 milliseconds
```

The filesystem under path '/data/bioproject.xml' is HEALTHY

This would explain the half-time only gain in speed.

Here are the times from assignment 1:

A screenshot of a terminal window titled "Downloads — ec2-user@ip-172-31-29-137:~ — ssh -i lwang73.pem ec2-user@e...". The terminal displays the output of a Hadoop job, showing various statistics such as "Spilled Records=3855505", "Shuffled Maps =2", "Failed Shuffles=0", "Merged Map outputs=2", "GC time elapsed (ms)=1086", "CPU time spent (ms)=41350", "Physical memory (bytes) snapshot=554782720", "Virtual memory (bytes) snapshot=6373744640", and "Total committed heap usage (bytes)=333918208". It also lists "Shuffle Errors" (BAD\_ID=0, CONNECTION=0, IO\_ERROR=0, WRONG\_LENGTH=0, WRONG\_MAP=0, WRONG\_REDUCE=0) and "File Input Format Counters" (Bytes Read=231153099) and "File Output Format Counters" (Bytes Written=20056175). At the bottom, it shows the execution times: "real 1m14.366s", "user 0m4.085s", and "sys 0m0.144s". The prompt is "[ec2-user@ip-172-31-29-137 ~]\$".

```
Spilled Records=3855505
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=1086
CPU time spent (ms)=41350
Physical memory (bytes) snapshot=554782720
Virtual memory (bytes) snapshot=6373744640
Total committed heap usage (bytes)=333918208

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=231153099
File Output Format Counters
  Bytes Written=20056175

real    1m14.366s
user    0m4.085s
sys     0m0.144s
[ec2-user@ip-172-31-29-137 ~]$
```

## Part 2: Hive

Run the following three (1.2, 1.3 and 2.1) queries in Hive and record the time they take to execute:

[http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/SSBM\\_queries.sql](http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/SSBM_queries.sql)

Create and load tables:

```
CREATE TABLE lineorder (
  lo_orderkey      INT,
  lo_linenumbers  INT,
  lo_custkey       INT,
  lo_partkey       INT,
```

```
lo_suppkey      INT,
lo_orderdate    INT,
lo_orderpriority VARCHAR(15),
lo_shippriority VARCHAR(1),
lo_quantity     INT,
lo_extendedprice INT,
lo_ordertotalprice INT,
lo_discount     INT,
lo_revenue      INT,
lo_supplycost   INT,
lo_tax          INT,
lo_commitdate   INT,
lo_shipmode     VARCHAR(10)
)
ROW FORMAT DELIMITED FIELDS
TERMINATED BY '|' STORED AS TEXTFILE;
```

```
LOAD DATA LOCAL INPATH '/home/ec2-user/lineorder.tbl' OVERWRITE INTO TABLE lineorder;
```

```
CREATE TABLE dwdate (
  d_datekey      INT,
  d_date         VARCHAR(19),
  d_dayofweek    VARCHAR(10),
  d_month        VARCHAR(10),
  d_year         INT,
  d_yearmonthnum INT,
  d_yearmonth    VARCHAR(8),
  d_daynuminweek INT,
  d_daynuminmonth INT,
  d_daynuminyear INT,
  d_monthnuminyear INT,
  d_weeknuminyear INT,
  d_sellingseason VARCHAR(13),
  d_lastdayinweekfl VARCHAR(1),
  d_lastdayinmonthfl VARCHAR(1),
  d_holidayfl    VARCHAR(1),
  d_weekdayfl     VARCHAR(1)
)
ROW FORMAT DELIMITED FIELDS
TERMINATED BY '|' STORED AS TEXTFILE;
```

```
LOAD DATA LOCAL INPATH '/home/ec2-user/dwdate.tbl' OVERWRITE INTO TABLE dwdate;
```

```
CREATE TABLE part (
  p_partkey INT,
  p_name    VARCHAR(22),
  p_mfgr    VARCHAR(6),
  p_category VARCHAR(7),
```

```
p_brand1  VARCHAR(9),
p_color   VARCHAR(11),
p_type    VARCHAR(25),
p_size    INT,
p_container VARCHAR(10)
)ROW FORMAT DELIMITED FIELDS
TERMINATED BY '|' STORED AS TEXTFILE;
```

```
LOAD DATA LOCAL INPATH '/home/ec2-user/part.tbl' OVERWRITE INTO TABLE part;
```

```
CREATE TABLE supplier (
  s_suppkey INT,
  s_name     VARCHAR(25),
  s_address  VARCHAR(25),
  s_city     VARCHAR(10),
  s_nation   VARCHAR(15),
  s_region   VARCHAR(12),
  s_phone    VARCHAR(15)
)ROW FORMAT DELIMITED FIELDS
TERMINATED BY '|' STORED AS TEXTFILE;
```

```
LOAD DATA LOCAL INPATH '/home/ec2-user/supplier.tbl' OVERWRITE INTO TABLE supplier;
```

## 1.2 Time to execute = 36.395 seconds

```
hive> select sum(lo_extendedprice) as revenue
> from lineorder, dwdate
> where lo_orderdate = d_datekey
> and d_yearmonth = 'Jan1993'
> and lo_discount between 5 and 6
> and lo_quantity between 25 and 35;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = ec2-user\_20190521013853\_cd4c2846-571d-49bd-9ed6-0da203aacb0e

Total jobs = 1

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in [jar:file:/home/ec2-user/apache-hive-2.0.1-bin/lib/log4j-slf4j-impl-2.4.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in [jar:file:/home/ec2-user/hadoop-2.6.4/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See [http://www.slf4j.org/codes.html#multiple\\_bindings](http://www.slf4j.org/codes.html#multiple_bindings) for an explanation.

SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Execution log at: /tmp/ec2-user/ec2-user\_20190521013853\_cd4c2846-571d-49bd-9ed6-0da203aacb0e.log

2019-05-21 01:38:59 Starting to launch local task to process map join; maximum memory = 477626368

## Lavinia Wang #1473704

2019-05-21 01:39:00 Dump the side-table for tag: 1 with group count: 31 into file: file:/tmp/ec2-user/5a9d12ba-e847-4d98-9d5f-3f8ad466871e/hive\_2019-05-21\_01-38-53\_599\_1551666118767653938-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile01--.hashtable  
2019-05-21 01:39:00 Uploaded 1 File to: file:/tmp/ec2-user/5a9d12ba-e847-4d98-9d5f-3f8ad466871e/hive\_2019-05-21\_01-38-53\_599\_1551666118767653938-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile01--.hashtable (945 bytes)  
2019-05-21 01:39:00 End of local task; Time Taken: 1.232 sec.  
Execution completed successfully  
MapredLocal task succeeded  
Launching Job 1 out of 1  
Number of reduce tasks determined at compile time: 1  
In order to change the average load for a reducer (in bytes):  
set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:  
set mapreduce.job.reduces=<number>  
Starting Job = job\_1558399145275\_0001, Tracking URL = http://ip-172-31-9-235.us-east-2.compute.internal:8088/proxy/application\_1558399145275\_0001/  
Kill Command = /home/ec2-user/hadoop-2.6.4/bin/hadoop job -kill job\_1558399145275\_0001  
Hadoop job information for Stage-2: number of mappers: 3; number of reducers: 1  
2019-05-21 01:39:08,201 Stage-2 map = 0%, reduce = 0%  
2019-05-21 01:39:21,323 Stage-2 map = 33%, reduce = 0%, Cumulative CPU 9.7 sec  
2019-05-21 01:39:24,590 Stage-2 map = 56%, reduce = 0%, Cumulative CPU 13.02 sec  
2019-05-21 01:39:25,749 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 13.72 sec  
2019-05-21 01:39:28,912 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 15.22 sec  
MapReduce Total cumulative CPU time: 15 seconds 220 msec  
Ended Job = job\_1558399145275\_0001  
MapReduce Jobs Launched:  
Stage-Stage-2: Map: 3 Reduce: 1 Cumulative CPU: 15.22 sec HDFS Read: 594368438 HDFS Write: 12  
SUCCESS  
Total MapReduce CPU Time Spent: 15 seconds 220 msec  
OK  
14215822897  
Time taken: 36.395 seconds, Fetched: 1 row(s)

### 1.3 Time to execute = 31.263 seconds

```
hive> select sum(lo_extendedprice) as revenue
> from lineorder, dwdate
> where lo_orderdate = d_datekey
> and d_weeknuminyear = 6 and d_year = 1994
> and lo_discount between 5 and 8
> and lo_quantity between 36 and 41;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = ec2-user\_20190521014210\_0b745812-8c9a-4c6d-b2b7-8dedb2fd9854

Total jobs = 1

## Lavinia Wang #1473704

SLF4J: Class path contains multiple SLF4J bindings.  
SLF4J: Found binding in [jar:file:/home/ec2-user/apache-hive-2.0.1-bin/lib/log4j-slf4j-impl-2.4.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]  
SLF4J: Found binding in [jar:file:/home/ec2-user/hadoop-2.6.4/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]  
SLF4J: See [http://www.slf4j.org/codes.html#multiple\\_bindings](http://www.slf4j.org/codes.html#multiple_bindings) for an explanation.  
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]  
Execution log at: /tmp/ec2-user/ec2-user\_20190521014210\_0b745812-8c9a-4c6d-b2b7-8dedb2fd9854.log  
2019-05-21 01:42:14 Starting to launch local task to process map join; maximum memory = 477626368  
2019-05-21 01:42:15 Dump the side-table for tag: 1 with group count: 7 into file: file:/tmp/ec2-user/5a9d12ba-e847-4d98-9d5f-3f8ad466871e/hive\_2019-05-21\_01-42-10\_401\_7461872208873411953-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile11--.hashtable  
2019-05-21 01:42:16 Uploaded 1 File to: file:/tmp/ec2-user/5a9d12ba-e847-4d98-9d5f-3f8ad466871e/hive\_2019-05-21\_01-42-10\_401\_7461872208873411953-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile11--.hashtable (414 bytes)  
2019-05-21 01:42:16 End of local task; Time Taken: 1.163 sec.  
Execution completed successfully  
MapredLocal task succeeded  
Launching Job 1 out of 1  
Number of reduce tasks determined at compile time: 1  
In order to change the average load for a reducer (in bytes):  
set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:  
set mapreduce.job.reduces=<number>  
Starting Job = job\_1558399145275\_0002, Tracking URL = [http://ip-172-31-9-235.us-east-2.compute.internal:8088/proxy/application\\_1558399145275\\_0002/](http://ip-172-31-9-235.us-east-2.compute.internal:8088/proxy/application_1558399145275_0002/)  
Kill Command = /home/ec2-user/hadoop-2.6.4/bin/hadoop job -kill job\_1558399145275\_0002  
Hadoop job information for Stage-2: number of mappers: 3; number of reducers: 1  
2019-05-21 01:42:21,508 Stage-2 map = 0%, reduce = 0%  
2019-05-21 01:42:33,313 Stage-2 map = 33%, reduce = 0%, Cumulative CPU 3.05 sec  
2019-05-21 01:42:34,425 Stage-2 map = 56%, reduce = 0%, Cumulative CPU 6.8 sec  
2019-05-21 01:42:35,470 Stage-2 map = 78%, reduce = 0%, Cumulative CPU 11.74 sec  
2019-05-21 01:42:36,496 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 12.86 sec  
2019-05-21 01:42:40,605 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 14.48 sec  
MapReduce Total cumulative CPU time: 14 seconds 480 msec  
Ended Job = job\_1558399145275\_0002  
MapReduce Jobs Launched:  
Stage-Stage-2: Map: 3 Reduce: 1 Cumulative CPU: 14.48 sec HDFS Read: 594368557 HDFS Write: 11  
SUCCESS  
Total MapReduce CPU Time Spent: 14 seconds 480 msec  
OK  
4435791464  
Time taken: 31.263 seconds, Fetched: 1 row(s)



## 2.1 Time to execute = 114.283 seconds

```
hive> select sum(lo_revenue), d_year, p_brand1
> from lineorder, dwdate, part, supplier
> where lo_orderdate = d_datekey
> and lo_partkey = p_partkey
> and lo_suppkey = s_suppkey
> and p_category = 'MFGR#12'
> and s_region = 'AMERICA'
> group by d_year, p_brand1
> order by d_year, p_brand1;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = ec2-user\_20190521014433\_b33e577a-c400-4c15-9266-3bf7cdf04727

Total jobs = 6

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in [jar:file:/home/ec2-user/apache-hive-2.0.1-bin/lib/log4j-slf4j-impl-2.4.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in [jar:file:/home/ec2-user/hadoop-2.6.4/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See [http://www.slf4j.org/codes.html#multiple\\_bindings](http://www.slf4j.org/codes.html#multiple_bindings) for an explanation.

SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Execution log at: /tmp/ec2-user/ec2-user\_20190521014433\_b33e577a-c400-4c15-9266-3bf7cdf04727.log

2019-05-21 01:44:38 Starting to launch local task to process map join; maximum memory = 477626368

2019-05-21 01:44:39 Dump the side-table for tag: 1 with group count: 2556 into file: file:/tmp/ec2-user/5a9d12ba-e847-4d98-9d5f-3f8ad466871e/hive\_2019-05-21\_01-44-33\_667\_5626930748980129596-1/-local-10014/HashTable-Stage-13/MapJoin-mapfile51--.hashtable

2019-05-21 01:44:39 Uploaded 1 File to: file:/tmp/ec2-user/5a9d12ba-e847-4d98-9d5f-3f8ad466871e/hive\_2019-05-21\_01-44-33\_667\_5626930748980129596-1/-local-10014/HashTable-Stage-13/MapJoin-mapfile51--.hashtable (67039 bytes)

2019-05-21 01:44:39 End of local task; Time Taken: 1.183 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 6

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job\_1558399145275\_0003, Tracking URL = [http://ip-172-31-9-235.us-east-2.compute.internal:8088/proxy/application\\_1558399145275\\_0003/](http://ip-172-31-9-235.us-east-2.compute.internal:8088/proxy/application_1558399145275_0003/)

Kill Command = /home/ec2-user/hadoop-2.6.4/bin/hadoop job -kill job\_1558399145275\_0003

Hadoop job information for Stage-13: number of mappers: 3; number of reducers: 0

2019-05-21 01:44:45,929 Stage-13 map = 0%, reduce = 0%

2019-05-21 01:45:01,911 Stage-13 map = 33%, reduce = 0%, Cumulative CPU 12.7 sec

2019-05-21 01:45:06,121 Stage-13 map = 50%, reduce = 0%, Cumulative CPU 21.24 sec

2019-05-21 01:45:08,260 Stage-13 map = 67%, reduce = 0%, Cumulative CPU 22.82 sec

2019-05-21 01:45:13,584 Stage-13 map = 83%, reduce = 0%, Cumulative CPU 28.51 sec

2019-05-21 01:45:14,607 Stage-13 map = 100%, reduce = 0%, Cumulative CPU 30.18 sec

MapReduce Total cumulative CPU time: 30 seconds 180 msec

Ended Job = job\_1558399145275\_0003

## Lavinia Wang #1473704

Stage-15 is selected by condition resolver.

Stage-16 is filtered out by condition resolver.

Stage-2 is filtered out by condition resolver.

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in [jar:file:/home/ec2-user/apache-hive-2.0.1-bin/lib/log4j-slf4j-impl-2.4.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in [jar:file:/home/ec2-user/hadoop-2.6.4/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See [http://www.slf4j.org/codes.html#multiple\\_bindings](http://www.slf4j.org/codes.html#multiple_bindings) for an explanation.

SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Execution log at: /tmp/ec2-user/ec2-user\_20190521014433\_b33e577a-c400-4c15-9266-3bf7cdf04727.log

2019-05-21 01:45:20 Starting to launch local task to process map join; maximum memory = 477626368

2019-05-21 01:45:21 Dump the side-table for tag: 1 with group count: 7883 into file: file:/tmp/ec2-user/5a9d12ba-e847-4d98-9d5f-3f8ad466871e/hive\_2019-05-21\_01-44-33\_667\_5626930748980129596-1/-local-10010/HashTable-Stage-10/MapJoin-mapfile31--.hashtable

2019-05-21 01:45:21 Uploaded 1 File to: file:/tmp/ec2-user/5a9d12ba-e847-4d98-9d5f-3f8ad466871e/hive\_2019-05-21\_01-44-33\_667\_5626930748980129596-1/-local-10010/HashTable-Stage-10/MapJoin-mapfile31--.hashtable (249337 bytes)

2019-05-21 01:45:21 End of local task; Time Taken: 1.761 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 3 out of 6

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job\_1558399145275\_0004, Tracking URL = [http://ip-172-31-9-235.us-east-2.compute.internal:8088/proxy/application\\_1558399145275\\_0004/](http://ip-172-31-9-235.us-east-2.compute.internal:8088/proxy/application_1558399145275_0004/)

Kill Command = /home/ec2-user/hadoop-2.6.4/bin/hadoop job -kill job\_1558399145275\_0004

Hadoop job information for Stage-10: number of mappers: 1; number of reducers: 0

2019-05-21 01:45:27,071 Stage-10 map = 0%, reduce = 0%

2019-05-21 01:45:40,419 Stage-10 map = 45%, reduce = 0%, Cumulative CPU 10.64 sec

2019-05-21 01:45:44,530 Stage-10 map = 100%, reduce = 0%, Cumulative CPU 14.75 sec

MapReduce Total cumulative CPU time: 14 seconds 750 msec

Ended Job = job\_1558399145275\_0004

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in [jar:file:/home/ec2-user/apache-hive-2.0.1-bin/lib/log4j-slf4j-impl-2.4.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in [jar:file:/home/ec2-user/hadoop-2.6.4/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See [http://www.slf4j.org/codes.html#multiple\\_bindings](http://www.slf4j.org/codes.html#multiple_bindings) for an explanation.

SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Execution log at: /tmp/ec2-user/ec2-user\_20190521014433\_b33e577a-c400-4c15-9266-3bf7cdf04727.log

2019-05-21 01:45:49 Starting to launch local task to process map join; maximum memory = 477626368

2019-05-21 01:45:51 Dump the side-table for tag: 1 with group count: 378 into file: file:/tmp/ec2-user/5a9d12ba-e847-4d98-9d5f-3f8ad466871e/hive\_2019-05-21\_01-44-33\_667\_5626930748980129596-1/-local-10008/HashTable-Stage-4/MapJoin-mapfile21--.hashtable

## Lavinia Wang #1473704

2019-05-21 01:45:51     Uploaded 1 File to: file:/tmp/ec2-user/5a9d12ba-e847-4d98-9d5f-3f8ad466871e/hive\_2019-05-21\_01-44-33\_667\_5626930748980129596-1/-local-10008/HashTable-Stage-4/MapJoin-mapfile21--.hashtable (7792 bytes)

2019-05-21 01:45:51     End of local task; Time Taken: 1.312 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 4 out of 6

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

    set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

    set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

    set mapreduce.job.reduces=<number>

Starting Job = job\_1558399145275\_0005, Tracking URL = [http://ip-172-31-9-235.us-east-2.compute.internal:8088/proxy/application\\_1558399145275\\_0005/](http://ip-172-31-9-235.us-east-2.compute.internal:8088/proxy/application_1558399145275_0005/)

Kill Command = /home/ec2-user/hadoop-2.6.4/bin/hadoop job -kill job\_1558399145275\_0005

Hadoop job information for Stage-4: number of mappers: 1; number of reducers: 1

2019-05-21 01:45:56,632 Stage-4 map = 0%, reduce = 0%

2019-05-21 01:46:02,806 Stage-4 map = 100%, reduce = 0%, Cumulative CPU 3.56 sec

2019-05-21 01:46:08,981 Stage-4 map = 100%, reduce = 100%, Cumulative CPU 5.2 sec

MapReduce Total cumulative CPU time: 5 seconds 200 msec

Ended Job = job\_1558399145275\_0005

Launching Job 5 out of 6

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

    set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

    set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

    set mapreduce.job.reduces=<number>

Starting Job = job\_1558399145275\_0006, Tracking URL = [http://ip-172-31-9-235.us-east-2.compute.internal:8088/proxy/application\\_1558399145275\\_0006/](http://ip-172-31-9-235.us-east-2.compute.internal:8088/proxy/application_1558399145275_0006/)

Kill Command = /home/ec2-user/hadoop-2.6.4/bin/hadoop job -kill job\_1558399145275\_0006

Hadoop job information for Stage-5: number of mappers: 1; number of reducers: 1

2019-05-21 01:46:16,561 Stage-5 map = 0%, reduce = 0%

2019-05-21 01:46:21,761 Stage-5 map = 100%, reduce = 0%, Cumulative CPU 1.02 sec

2019-05-21 01:46:26,896 Stage-5 map = 100%, reduce = 100%, Cumulative CPU 2.3 sec

MapReduce Total cumulative CPU time: 2 seconds 300 msec

Ended Job = job\_1558399145275\_0006

MapReduce Jobs Launched:

Stage-Stage-13: Map: 3   Cumulative CPU: 30.18 sec   HDFS Read: 594359050 HDFS Write: 184733291  
SUCCESS

Stage-Stage-10: Map: 1   Cumulative CPU: 14.75 sec   HDFS Read: 184739547 HDFS Write: 8750735  
SUCCESS

Stage-Stage-4: Map: 1 Reduce: 1   Cumulative CPU: 5.2 sec   HDFS Read: 8762980 HDFS Write: 9913  
SUCCESS

## Lavinia Wang #1473704

Stage-Stage-5: Map: 1 Reduce: 1 Cumulative CPU: 2.3 sec HDFS Read: 15707 HDFS Write: 6937  
SUCCESS

Total MapReduce CPU Time Spent: 52 seconds 430 msec

OK

567838207 1992 MFGR#121  
610663790 1992 MFGR#1210  
550769662 1992 MFGR#1211  
649205856 1992 MFGR#1212  
624031241 1992 MFGR#1213  
670488468 1992 MFGR#1214  
633152470 1992 MFGR#1215  
674846781 1992 MFGR#1216  
675093435 1992 MFGR#1217  
600202070 1992 MFGR#1218  
538043594 1992 MFGR#1219  
655326672 1992 MFGR#122  
540262882 1992 MFGR#1220  
556120633 1992 MFGR#1221  
590762777 1992 MFGR#1222  
535448651 1992 MFGR#1223  
703752611 1992 MFGR#1224  
570832868 1992 MFGR#1225  
614061593 1992 MFGR#1226  
581759388 1992 MFGR#1227  
644642592 1992 MFGR#1228  
640858430 1992 MFGR#1229  
789755835 1992 MFGR#123  
468535087 1992 MFGR#1230  
592436656 1992 MFGR#1231  
664275152 1992 MFGR#1232  
613885100 1992 MFGR#1233  
667399281 1992 MFGR#1234  
640290070 1992 MFGR#1235  
501892561 1992 MFGR#1236  
591481503 1992 MFGR#1237  
477423770 1992 MFGR#1238  
638259374 1992 MFGR#1239  
572354196 1992 MFGR#124  
740479248 1992 MFGR#1240  
478777095 1992 MFGR#125  
592174616 1992 MFGR#126  
706151632 1992 MFGR#127  
542306646 1992 MFGR#128  
581987352 1992 MFGR#129  
823087702 1993 MFGR#121  
648160706 1993 MFGR#1210  
634743898 1993 MFGR#1211  
785639283 1993 MFGR#1212

Lavinia Wang #1473704

638255029 1993 MFGR#1213  
616837237 1993 MFGR#1214  
634687975 1993 MFGR#1215  
638353900 1993 MFGR#1216  
663372951 1993 MFGR#1217  
683985855 1993 MFGR#1218  
646950033 1993 MFGR#1219  
622532984 1993 MFGR#122  
530830127 1993 MFGR#1220  
543346337 1993 MFGR#1221  
756921203 1993 MFGR#1222  
533544350 1993 MFGR#1223  
915916085 1993 MFGR#1224  
473007381 1993 MFGR#1225  
739036124 1993 MFGR#1226  
592178887 1993 MFGR#1227  
583507058 1993 MFGR#1228  
617453491 1993 MFGR#1229  
637863868 1993 MFGR#123  
625534310 1993 MFGR#1230  
580327635 1993 MFGR#1231  
697373098 1993 MFGR#1232  
515571416 1993 MFGR#1233  
651935758 1993 MFGR#1234  
575779480 1993 MFGR#1235  
591878667 1993 MFGR#1236  
609618576 1993 MFGR#1237  
444614010 1993 MFGR#1238  
595256327 1993 MFGR#1239  
660586237 1993 MFGR#124  
788730059 1993 MFGR#1240  
616224539 1993 MFGR#125  
617126754 1993 MFGR#126  
654438324 1993 MFGR#127  
731657001 1993 MFGR#128  
548048395 1993 MFGR#129  
564405648 1994 MFGR#121  
645404849 1994 MFGR#1210  
631620635 1994 MFGR#1211  
568332348 1994 MFGR#1212  
678785857 1994 MFGR#1213  
534002330 1994 MFGR#1214  
654400242 1994 MFGR#1215  
558646341 1994 MFGR#1216  
687845641 1994 MFGR#1217  
546674347 1994 MFGR#1218  
567272942 1994 MFGR#1219  
659884062 1994 MFGR#122

Lavinia Wang #1473704

562582172	1994	MFGR#1220
598618997	1994	MFGR#1221
601016441	1994	MFGR#1222
555134404	1994	MFGR#1223
737422302	1994	MFGR#1224
570745955	1994	MFGR#1225
746302245	1994	MFGR#1226
651707481	1994	MFGR#1227
573693547	1994	MFGR#1228
647918373	1994	MFGR#1229
580449592	1994	MFGR#123
493270412	1994	MFGR#1230
603546148	1994	MFGR#1231
719865331	1994	MFGR#1232
638982238	1994	MFGR#1233
743247677	1994	MFGR#1234
598680959	1994	MFGR#1235
615726097	1994	MFGR#1236
542569815	1994	MFGR#1237
573510781	1994	MFGR#1238
579855853	1994	MFGR#1239
684573322	1994	MFGR#124
873735737	1994	MFGR#1240
560488304	1994	MFGR#125
657036514	1994	MFGR#126
622571183	1994	MFGR#127
586845664	1994	MFGR#128
534541525	1994	MFGR#129
706469511	1995	MFGR#121
602892803	1995	MFGR#1210
645166092	1995	MFGR#1211
613289283	1995	MFGR#1212
599586479	1995	MFGR#1213
562570804	1995	MFGR#1214
672528755	1995	MFGR#1215
669000972	1995	MFGR#1216
725362449	1995	MFGR#1217
657026635	1995	MFGR#1218
519659003	1995	MFGR#1219
724727741	1995	MFGR#122
517956131	1995	MFGR#1220
635741351	1995	MFGR#1221
564368410	1995	MFGR#1222
600665149	1995	MFGR#1223
762700351	1995	MFGR#1224
671669586	1995	MFGR#1225
572568748	1995	MFGR#1226
530361300	1995	MFGR#1227

Lavinia Wang #1473704

633357085 1995 MFGR#1228  
547960244 1995 MFGR#1229  
660711077 1995 MFGR#123  
602735858 1995 MFGR#1230  
499852146 1995 MFGR#1231  
715300753 1995 MFGR#1232  
557149571 1995 MFGR#1233  
710023059 1995 MFGR#1234  
622425239 1995 MFGR#1235  
634565501 1995 MFGR#1236  
572847270 1995 MFGR#1237  
549318912 1995 MFGR#1238  
593851712 1995 MFGR#1239  
585421815 1995 MFGR#124  
707207888 1995 MFGR#1240  
538246872 1995 MFGR#125  
605799021 1995 MFGR#126  
665978112 1995 MFGR#127  
646960956 1995 MFGR#128  
508749401 1995 MFGR#129  
523879145 1996 MFGR#121  
643645053 1996 MFGR#1210  
595065339 1996 MFGR#1211  
674626440 1996 MFGR#1212  
496297087 1996 MFGR#1213  
583249505 1996 MFGR#1214  
702184857 1996 MFGR#1215  
601809334 1996 MFGR#1216  
704898387 1996 MFGR#1217  
528843086 1996 MFGR#1218  
586246330 1996 MFGR#1219  
712110492 1996 MFGR#122  
518444215 1996 MFGR#1220  
499319414 1996 MFGR#1221  
679469356 1996 MFGR#1222  
628762754 1996 MFGR#1223  
724844856 1996 MFGR#1224  
660620587 1996 MFGR#1225  
667674729 1996 MFGR#1226  
483838085 1996 MFGR#1227  
609855391 1996 MFGR#1228  
658959557 1996 MFGR#1229  
566217852 1996 MFGR#123  
528879998 1996 MFGR#1230  
589481194 1996 MFGR#1231  
702805896 1996 MFGR#1232  
663679947 1996 MFGR#1233  
571149450 1996 MFGR#1234



Lavinia Wang #1473704

478648074 1996 MFGR#1235  
568249365 1996 MFGR#1236  
592616167 1996 MFGR#1237  
466676148 1996 MFGR#1238  
670693719 1996 MFGR#1239  
560667719 1996 MFGR#124  
821167950 1996 MFGR#1240  
476864333 1996 MFGR#125  
558030884 1996 MFGR#126  
635873891 1996 MFGR#127  
551010618 1996 MFGR#128  
560570630 1996 MFGR#129  
587013207 1997 MFGR#121  
616287892 1997 MFGR#1210  
548588761 1997 MFGR#1211  
589593892 1997 MFGR#1212  
424306670 1997 MFGR#1213  
511971910 1997 MFGR#1214  
631772246 1997 MFGR#1215  
692135140 1997 MFGR#1216  
777994957 1997 MFGR#1217  
707053720 1997 MFGR#1218  
561169527 1997 MFGR#1219  
664916245 1997 MFGR#122  
594466157 1997 MFGR#1220  
588848171 1997 MFGR#1221  
528988960 1997 MFGR#1222  
537098211 1997 MFGR#1223  
674763166 1997 MFGR#1224  
450402292 1997 MFGR#1225  
701360722 1997 MFGR#1226  
506011570 1997 MFGR#1227  
585578737 1997 MFGR#1228  
622744016 1997 MFGR#1229  
646503168 1997 MFGR#123  
571800941 1997 MFGR#1230  
502601790 1997 MFGR#1231  
677924656 1997 MFGR#1232  
534455976 1997 MFGR#1233  
714934715 1997 MFGR#1234  
767151420 1997 MFGR#1235  
618877179 1997 MFGR#1236  
639638057 1997 MFGR#1237  
401953419 1997 MFGR#1238  
610756714 1997 MFGR#1239  
543248087 1997 MFGR#124  
675132692 1997 MFGR#1240  
479099365 1997 MFGR#125

## Lavinia Wang #1473704

570696568	1997	MFGR#126
583074592	1997	MFGR#127
695133104	1997	MFGR#128
655638776	1997	MFGR#129
344575925	1998	MFGR#121
417152416	1998	MFGR#1210
317068168	1998	MFGR#1211
374341516	1998	MFGR#1212
332740903	1998	MFGR#1213
304873002	1998	MFGR#1214
366101132	1998	MFGR#1215
379133898	1998	MFGR#1216
359508497	1998	MFGR#1217
320623334	1998	MFGR#1218
346182862	1998	MFGR#1219
312440027	1998	MFGR#122
348123961	1998	MFGR#1220
339845398	1998	MFGR#1221
355416161	1998	MFGR#1222
344889822	1998	MFGR#1223
396906691	1998	MFGR#1224
290208878	1998	MFGR#1225
419415707	1998	MFGR#1226
358466340	1998	MFGR#1227
251549955	1998	MFGR#1228
383138860	1998	MFGR#1229
296330561	1998	MFGR#123
437181243	1998	MFGR#1230
398944492	1998	MFGR#1231
424062455	1998	MFGR#1232
406967188	1998	MFGR#1233
428867240	1998	MFGR#1234
352277781	1998	MFGR#1235
361827086	1998	MFGR#1236
341618569	1998	MFGR#1237
244739231	1998	MFGR#1238
414151803	1998	MFGR#1239
330082371	1998	MFGR#124
415312453	1998	MFGR#1240
360289624	1998	MFGR#125
341657580	1998	MFGR#126
377507061	1998	MFGR#127
361416497	1998	MFGR#128
318769573	1998	MFGR#129

Time taken: 114.283 seconds, Fetched: 280 row(s)

Perform the following transform operation using SELECT TRANSFORM on the customer table by creating a new table. Your new target table should have only three columns, c\_custkey (no changes), c\_address, and c\_city.

## Lavinia Wang #1473704

For the c\_address column, shorten it to 8 characters (i.e., if the value is longer, remove extra characters, but otherwise keep it as-is). For c\_city, add a space and a # to indicate the digit at the end (e.g., UNITED KI2 => UNITED KI #2, or INDONESIA4 => INDONESIA #4). Make sure to modify the columns of the target table accordingly (since you are introducing longer columns).

### customer.py

```
#!/usr/bin/python
import sys

for line in sys.stdin:
    line = line.strip().split(',')
    line[2] = line[2][:8]
    line[3] = line[3][:len(line[3])-1] + ' #' + line[3][len(line[3])-1]
    print '\t'.join(line)
```

### create and load customer table

```
CREATE TABLE customer (
  c_custkey INT,
  c_name VARCHAR(25),
  c_address VARCHAR(25),
  c_city VARCHAR(10),
  c_nation VARCHAR(15),
  c_region VARCHAR(12),
  c_phone VARCHAR(15),
  c_mktsegment VARCHAR(10)
)
ROW FORMAT DELIMITED FIELDS
TERMINATED BY '|' STORED AS TEXTFILE;
```

```
LOAD DATA LOCAL INPATH '/home/ec2-user/customer.tbl' OVERWRITE INTO TABLE customer;
```

### Create customer2 table

```
create table customer2 (
  c_custkey INT,
  c_name VARCHAR(25),
  c_address VARCHAR(8),
  c_city VARCHAR(12),
  c_nation VARCHAR(15),
  c_region VARCHAR(12),
  c_phone VARCHAR(15),
  c_mktsegment VARCHAR(10)
)
ROW FORMAT DELIMITED FIELDS
TERMINATED BY '\t' STORED AS TEXTFILE;
```

### Select transform

## Lavinia Wang #1473704

```
INSERT OVERWRITE TABLE customer2 SELECT TRANSFORM (c_custkey, c_name, c_address, c_city, c_nation,
c_region, c_phone, c_mktsegment) USING 'python customer.py' AS (c_custkey, c_name, c_address, c_city,
c_nation, c_region, c_phone, c_mktsegment) FROM customer;
```

### test query:

```
hive> select * from customer2 limit 5;
```

OK

```
1  Customer#000000001  j5JsirBM  MOROCCO  #0  MOROCCO AFRICA  25-989-741-2988 BUILDING
2  Customer#000000002  487LW1do  JORDAN  #1  JORDAN MIDDLE EAST  23-768-687-3665 AUTOMOBILE
3  Customer#000000003  fkRGN8n ARGENTINA #7  ARGENTINA  AMERICA 11-719-748-3364 AUTOMOBILE
4  Customer#000000004  4u58h f EGYPT  #4  EGYPT  MIDDLE EAST  14-128-190-5944 MACHINERY
5  Customer#000000005  hwBtxkoB  CANADA  #5  CANADA AMERICA 13-750-942-6364 HOUSEHOLD
```

Time taken: 0.051 seconds, Fetched: 5 row(s)

## Part 3: Pig

Convert and load the data into Pig, implementing only queries 0.1, 0.2, 0.3. Do not implement all queries.

Check disk storage space in HDFS, if your disk usage is over 90% Pig may hang without an error or a warning.

One easy way to time Pig is as follows: put your sequence of pig commands into a text file and then run, from command line in pig directory (e.g., [ec2-user@ip-172-31-6-39 pig-0.15.0]\$), **bin/pig -f pig\_script.pig** (which will inform you how long the pig script took to run).

### 0.1

```
lineorder = LOAD '/user/ec2-user/lineorder.tbl' USING PigStorage('|')
AS (lo_orderkey:int,
   lo_linenummer:int,
   lo_custkey:int,
   lo_partkey:int,
   lo_suppkey:int,
   lo_orderdate:int,
   lo_orderpriority:chararray,
   lo_shippriority:chararray,
   lo_quantity:int,
   lo_extendedprice:int,
   lo_ordertotalprice:int,
   lo_discount:int,
   lo_revenue:int,
```

```
lo_supplycost:int,  
lo_tax:int,  
lo_commitdate:int,  
lo_shipmode::chararray  
);
```

```
grouped = group lineorder all;  
avg = FOREACH grouped GENERATE AVG(lineorder.lo_revenue);  
DUMP avg;
```

Answer = 3634300.709514323

Execution time of 3 minutes, 34 seconds, 11 millisecond

## 0.2

```
grouped = group lineorder by lo_discount;  
counted = FOREACH grouped GENERATE group as lo_discount, COUNT(lineorder.lo_extendedprice) as  
cnt;  
DUMP counted;
```

2019-05-21 05:20:36,325 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil  
- Total input paths to process : 1

(0,544886)  
(1,545834)  
(2,546173)  
(3,545293)  
(4,545545)  
(5,546395)  
(6,544970)  
(7,546192)  
(8,544803)  
(9,545309)  
(10,545815)

2019-05-21 05:20:36,402 [main] INFO org.apache.pig.Main - Pig script completed in 3 minutes, 47  
seconds and 31 milliseconds (227031ms)

## 0.3

```
filtered = FILTER lineorder BY lo_discount < 3;  
grouped = GROUP filtered BY lo_quantity;  
summed = FOREACH grouped GENERATE group as lo_quantity, SUM(filtered.lo_revenue) as rev;  
DUMP summed;
```

2019-05-21 06:04:22,932 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil  
- Total input paths to process : 1  
(1,4879019020)

Lavinia Wang #1473704

(2,9644127315)  
(3,14575887127)  
(4,19360189865)  
(5,24073923574)  
(6,29125189531)  
(7,33982891466)  
(8,38671565454)  
(9,43381602619)  
(10,48619780003)  
(11,53159489411)  
(12,58264291629)  
(13,62920595763)  
(14,67451818069)  
(15,73414895616)  
(16,78360133885)  
(17,82320521791)  
(18,86995495639)  
(19,93016668045)  
(20,97258062753)  
(21,100684344044)  
(22,107454001560)  
(23,112984674102)  
(24,116527702603)  
(25,123160894092)  
(26,126451771059)  
(27,132113291310)  
(28,135413154368)  
(29,141357789043)  
(30,145181046794)  
(31,149937771539)  
(32,157770330201)  
(33,161774040572)  
(34,164150363629)  
(35,170173151151)  
(36,175712858188)  
(37,178733976488)  
(38,186428562667)  
(39,187696104837)  
(40,196345645204)  
(41,199250645070)  
(42,204966410590)  
(43,209016181876)  
(44,213245636104)  
(45,217565230742)  
(46,223784510215)  
(47,229077142619)  
(48,234125822088)  
(49,236641410613)

(50,243791122644)

2019-05-21 06:04:22,987 [main] INFO org.apache.pig.Main - Pig script completed in 3 minutes, 17 seconds and 45 milliseconds (197045 ms)

## Part 4: Hadoop Streaming

Implement query **0.3** using Hadoop streaming with python. You don't need to implement other queries.

```
--Q0.3
SELECT lo_quantity, SUM(lo_revenue)
FROM lineorder
WHERE lo_discount < 3
GROUP BY lo_quantity;
```

**Set up commands to move lineorder.tbl into Hadoop directory.**

```
[ec2-user@ip-172-31-9-235 ~]$ hadoop fs -mkdir /user/ec2-user/ssbm
[ec2-user@ip-172-31-9-235 ~]$ hadoop fs -put lineorder.tbl ssbm
[ec2-user@ip-172-31-9-235 ~]$ hadoop fs -ls ssbm
Found 1 items
-rw-r--r--  2 ec2-user supergroup  594313001 2019-05-21 03:01 ssbm/lineorder.tbl
```

### myMapper.0.3.py

```
#!/usr/bin/python

import sys

for line in sys.stdin:
    line = line.strip()
    vals = line.split("|")

    discount = int(vals[11])
    quantity = vals[8]
    revenue = vals[12]

    if discount < 3:
        print "%s\t%s" % (quantity, revenue)
```

### myReducer.0.3.py

```
#!/usr/bin/python
import sys

curr_id = None
curr_tot = 0
id = None
```



## Lavinia Wang #1473704

```
# The input comes from standard input (line by line)
for line in sys.stdin:
    line = line.strip()
    # parse the line and split it by '\t'
    ln = line.split('\t')
    # grab the key
    id = ln[0]
    # grab the value (int)
    val = int(ln[1])

    if curr_id == id:
        curr_tot += val
    else:
        if curr_id: # output the sum, single key completed
            print '%s\t%d' % (curr_id, curr_tot)

            curr_id = id
            curr_tot = val

# output the last key
if curr_id == id:
    print '%s\t%d' % (curr_id, curr_tot)
```

### command

```
[ec2-user@ip-172-31-9-235~]$ hadoop jar hadoop-streaming-2.6.4.jar -input
/user/ec2-user/ssbm -output /data/output16 -mapper myMapper.0.3.py -reducer
myReducer.0.3.py -file myReducer.0.3.py -file myMapper.0.3.py
```

### Results

```
[ec2-user@ip-172-31-9-235~]$ hadoop fs -cat /data/output16/part-00000
1      4879019020
10     48619780003
11     53159489411
12     58264291629
13     62920595763
14     67451818069
15     73414895616
16     78360133885
17     82320521791
18     86995495639
19     93016668045
2      9644127315
20     97258062753
21     100684344044
22     107454001560
23     112984674102
24     116527702603
25     123160894092
26     126451771059
27     132113291310
28     135413154368
29     141357789043
3      14575887127
30     145181046794
```

## Lavinia Wang #1473704

31	149937771539
32	157770330201
33	161774040572
34	164150363629
35	170173151151
36	175712858188
37	178733976488
38	186428562667
39	187696104837
4	19360189865
40	196345645204
41	199250645070
42	204966410590
43	209016181876
44	213245636104
45	217565230742
46	223784510215
47	229077142619
48	234125822088
49	236641410613
5	24073923574
50	243791122644
6	29125189531
7	33982891466
8	38671565454
9	43381602619

NOTE: You may implement this part in Java if you prefer.

Submit a single document containing your written answers. Be sure that this document contains your name and “CSC 555 Project Phase 1” at the top.